Package 'tidytext'

April 10, 2021

Type Package

Title Text Mining using 'dplyr', 'ggplot2', and Other Tidy Tools

Version 0.3.1

Description Using tidy data principles can make many text mining tasks easier,
more effective, and consistent with tools already in wide use. Much of the
infrastructure needed for text mining with tidy data frames already exists
in packages like 'dplyr', 'broom', 'tidyr', and 'ggplot2'. In this package,
we provide functions and supporting data sets to allow conversion of text
to and from tidy formats, and to switch seamlessly between tidy tools and
existing text mining packages.

License MIT + file LICENSE

Encoding UTF-8

LazyData TRUE

URL https://github.com/juliasilge/tidytext

BugReports https://github.com/juliasilge/tidytext/issues

RoxygenNote 7.1.1.9001

Depends R (>= 2.10)

Imports rlang (>= 0.4.10), tibble, dplyr, stringr, hunspell, generics,
lifecycle, Matrix, tokenizers, janeaustenr, purrr (>= 0.1.1),
methods, vctrs

Suggests readr, tidyr, XML, tm, quanteda, knitr, rmarkdown, ggplot2,
reshape2, wordcloud, topicmodels, NLP, scales, testthat (>=
2.1.0), vdiffr, mallet, stm, data.table, broom, textdata,
stopwords, covr

VignetteBuilder knitr

NeedsCompilation no

Author Gabriela De Queiroz [ctb],
Colin Fay [ctb] (<https://orcid.org/0000-0001-7343-1846>),
Emil Hvitfeldt [ctb],
Os Keyes [ctb] (<https://orcid.org/0000-0001-5196-609X>),
Kanishka Misra [ctb],

Tim Mastny [ctb],
Jeff Erickson [ctb],
David Robinson [aut],
Julia Silge [aut, cre] (<https://orcid.org/0000-0002-3671-836X>)

Maintainer Julia Silge <julia.silge@gmail.com>

R **topics documented:**

| bind_tf_idf | *Bind the term frequency and inverse document frequency of a tidy text dataset to the dataset* |

Description

Calculate and bind the term frequency and inverse document frequency of a tidy text dataset, along with the product, tf-idf, to the dataset. Each of these values are added as columns. This function supports non-standard evaluation through the tidyeval framework.

Usage

```
bind_tf_idf(tbl, term, document, n)
```

Arguments

tbl	A tidy text dataset with one-row-per-term-per-document
term	Column containing terms as string or symbol
document	Column containing document IDs as string or symbol
n	Column containing document-term counts as string or symbol

Details

The arguments term, document, and n are passed by expression and support [quasiquotation](); you can unquote strings and symbols.

If the dataset is grouped, the groups are ignored but are retained.

The dataset must have exactly one row per document-term combination for this to work.

Examples

```
library(dplyr)
library(janeaustenr)

book_words <- austen_books() %>%
  unnest_tokens(word, text) %>%
  count(book, word, sort = TRUE)

book_words

# find the words most distinctive to each document
book_words %>%
  bind_tf_idf(word, book, n) %>%
  arrange(desc(tf_idf))
```

---

cast_sparse	*Create a sparse matrix from row names, column names, and values in a table.*

Description

This function supports non-standard evaluation through the tidyeval framework.

Usage

```
cast_sparse(data, row, column, value, ...)
```

Arguments

data	A tbl
row	Column name to use as row names in sparse matrix, as string or symbol
column	Column name to use as column names in sparse matrix, as string or symbol
value	Column name to use as sparse matrix values (default 1) as string or symbol
...	Extra arguments to pass on to [sparseMatrix](#)

Details

Note that cast_sparse ignores groups in a grouped tbl_df. The arguments `row`, `column`, and `value` are passed by expression and support [quasiquotation](#); you can unquote strings and symbols.

Value

A sparse Matrix object, with one row for each unique value in the `row` column, one column for each unique value in the `column` column, and with as many non-zero values as there are rows in `data`.

Examples

```
dat <- data.frame(a = c("row1", "row1", "row2", "row2", "row2"),
                  b = c("col1", "col2", "col1", "col3", "col4"),
                  val = 1:5)

cast_sparse(dat, a, b)

cast_sparse(dat, a, b, val)
```

---

cast_tdm	*Casting a data frame to a DocumentTermMatrix, TermDocumentMatrix, or dfm*

Description

This turns a "tidy" one-term-per-document-per-row data frame into a DocumentTermMatrix or TermDocumentMatrix from the tm package, or a dfm from the quanteda package. These functions support non-standard evaluation through the tidyeval framework. Groups are ignored.

Usage

```
cast_tdm(data, term, document, value, weighting = tm::weightTf, ...)

cast_dtm(data, document, term, value, weighting = tm::weightTf, ...)

cast_dfm(data, document, term, value, ...)
```

Arguments

data	Table with one-term-per-document-per-row
term	Column containing terms as string or symbol
document	Column containing document IDs as string or symbol
value	Column containing values as string or symbol
weighting	The weighting function for the DTM/TDM (default is term-frequency, effectively unweighted)
...	Extra arguments passed on to `sparseMatrix`

Details

The arguments `term`, `document`, and `value` are passed by expression and support [quasiquotation](); you can unquote strings and symbols.

---

corpus_tidiers	*Tidiers for a corpus object from the quanteda package*

Description

Tidy a corpus object from the quanteda package. `tidy` returns a tbl_df with one-row-per-document, with a `text` column containing the document's text, and one column for each document-level metadata. `glance` returns a one-row tbl_df with corpus-level metadata, such as source and created. For Corpus objects from the tm package, see `tidy.Corpus`.

Usage

```
## S3 method for class 'corpus'
tidy(x, ...)

## S3 method for class 'corpus'
glance(x, ...)
```

Arguments

x	A Corpus object, such as a VCorpus or PCorpus
...	Extra arguments, not used

Details

For the most part, the tidy output is equivalent to the "documents" data frame in the corpus object, except that it is converted to a tbl_df, and texts column is renamed to text to be consistent with other uses in tidytext.

Similarly, the glance output is simply the "metadata" object, with NULL fields removed and turned into a one-row tbl_df.

Examples

```
if (requireNamespace("quanteda", quietly = TRUE)) {
 data("data_corpus_inaugural", package = "quanteda")

 data_corpus_inaugural

 tidy(data_corpus_inaugural)
}
```

---

dictionary_tidiers   *Tidy dictionary objects from the quanteda package*

Description

Tidy dictionary objects from the quanteda package

Usage

```
## S3 method for class 'dictionary2'
tidy(x, regex = FALSE, ...)
```

Arguments

x	A dictionary object
regex	Whether to turn dictionary items from a glob to a regex
...	Extra arguments, not used

Value

A data frame with two columns: category and word.

---

get_sentiments	*Get a tidy data frame of a single sentiment lexicon*

Description

Get specific sentiment lexicons in a tidy format, with one row per word, in a form that can be joined with a one-word-per-row dataset. The "bing" option comes from the included [sentiments](#) data frame, and others call the relevant function in the **textdata** package.

Usage

```
get_sentiments(lexicon = c("bing", "afinn", "loughran", "nrc"))
```

Arguments

lexicon	The sentiment lexicon to retrieve; either "afinn", "bing", "nrc", or "loughran"

Value

A tbl_df with a word column, and either a sentiment column (if lexicon is not "afinn") or a numeric value column (if lexicon is "afinn").

Examples

```
library(dplyr)

get_sentiments("bing")

## Not run:
get_sentiments("afinn")
get_sentiments("nrc")

## End(Not run)
```

---

get_stopwords   *Get a tidy data frame of a single stopword lexicon*

Description

Get a specific stop word lexicon via the **stopwords** package's [stopwords](#) function, in a tidy format with one word per row.

Usage

```
get_stopwords(language = "en", source = "snowball")
```

Arguments

language   The language of the stopword lexicon specified as a two-letter ISO code, such as
              "es", "de", or "fr". Default is "en" for English. Use [stopwords_getlanguages](#)
              from **stopwords** to see available languages.

source   The source of the stopword lexicon specified. Default is "snowball". Use
              [stopwords_getsources](#) from **stopwords** to see available sources.

Value

A tibble with two columns, word and lexicon. The parameter lexicon is "quanteda" in this case.

Examples

```
library(dplyr)
get_stopwords()
get_stopwords(source = "smart")
get_stopwords("es", "snowball")
get_stopwords("ru", "snowball")
```

---

lda_tidiers   *Tidiers for LDA and CTM objects from the topicmodels package*

Description

Tidy the results of a Latent Dirichlet Allocation or Correlated Topic Model.

Usage

```
## S3 method for class 'LDA'
tidy(x, matrix = c("beta", "gamma"), log = FALSE, ...)

## S3 method for class 'CTM'
tidy(x, matrix = c("beta", "gamma"), log = FALSE, ...)

## S3 method for class 'LDA'
augment(x, data, ...)

## S3 method for class 'CTM'
augment(x, data, ...)

## S3 method for class 'LDA'
glance(x, ...)

## S3 method for class 'CTM'
glance(x, ...)
```

Arguments

x	An LDA or CTM (or LDA_VEM/CTA_VEM) object from the topicmodels package
matrix	Whether to tidy the beta (per-term-per-topic, default) or gamma (per-document-per-topic) matrix
log	Whether beta/gamma should be on a log scale, default FALSE
...	Extra arguments, not used
data	For augment, the data given to the LDA or CTM function, either as a DocumentTermMatrix or as a tidied table with "document" and "term" columns

Value

`tidy` returns a tidied version of either the beta or gamma matrix.

If `matrix == "beta"` (default), returns a table with one row per topic and term, with columns

topic Topic, as an integer

term Term

beta Probability of a term generated from a topic according to the multinomial model

If `matrix == "gamma"`, returns a table with one row per topic and document, with columns

topic Topic, as an integer

document Document name or ID

gamma Probability of topic given document

augment returns a table with one row per original document-term pair, such as is returned by tdm_tidiers:

document Name of document (if present), or index

term Term

.topic Topic assignment

If the `data` argument is provided, any columns in the original data are included, combined based on the `document` and `term` columns.

`glance` always returns a one-row table, with columns

iter Number of iterations used

terms Number of terms in the model

alpha If an LDA_VEM, the parameter of the Dirichlet distribution for topics over documents

Examples

```
if (requireNamespace("topicmodels", quietly = TRUE)) {
  set.seed(2016)
  library(dplyr)
  library(topicmodels)

  data("AssociatedPress", package = "topicmodels")
  ap <- AssociatedPress[1:100, ]
  lda <- LDA(ap, control = list(alpha = 0.1), k = 4)

  # get term distribution within each topic
  td_lda <- tidy(lda)
  td_lda

  library(ggplot2)

  # visualize the top terms within each topic
  td_lda_filtered <- td_lda %>%
    filter(beta > .004) %>%
    mutate(term = reorder(term, beta))

  ggplot(td_lda_filtered, aes(term, beta)) +
    geom_bar(stat = "identity") +
    facet_wrap(~ topic, scales = "free") +
    theme(axis.text.x = element_text(angle = 90, size = 15))

  # get classification of each document
  td_lda_docs <- tidy(lda, matrix = "gamma")
  td_lda_docs

  doc_classes <- td_lda_docs %>%
    group_by(document) %>%
    top_n(1) %>%
    ungroup()

  doc_classes
```

```
  # which were we most uncertain about?
  doc_classes %>%
    arrange(gamma)
}
```

---

mallet_tidiers   *Tidiers for Latent Dirichlet Allocation models from the mallet package*

Description

Tidy LDA models fit by the mallet package, which wraps the Mallet topic modeling package in Java. The arguments and return values are similar to `lda_tidiers`.

Usage

```
## S3 method for class 'jobjRef'
tidy(
  x,
  matrix = c("beta", "gamma"),
  log = FALSE,
  normalized = TRUE,
  smoothed = TRUE,
  ...
)

## S3 method for class 'jobjRef'
augment(x, data, ...)
```

Arguments

x	A jobjRef object, of type RTopicModel, such as created by `MalletLDA`.
matrix	Whether to tidy the beta (per-term-per-topic, default) or gamma (per-document-per-topic) matrix.
log	Whether beta/gamma should be on a log scale, default FALSE
normalized	If true (default), normalize so that each document or word sums to one across the topics. If false, values will be integers representing the actual number of word-topic or document-topic assignments.
smoothed	If true (default), add the smoothing parameter to each to avoid any values being zero. This smoothing parameter is initialized as `alpha.sum` in `MalletLDA`.
...	Extra arguments, not used
data	For augment, the data given to the LDA function, either as a DocumentTermMatrix or as a tidied table with "document" and "term" columns.

Details

Note that the LDA models from `MalletLDA` are technically a special case of S4 objects with class jobjRef. These are thus implemented as jobjRef tidiers, with a check for whether the `toString` output is as expected.

Value

augment must be provided a data argument containing one row per original document-term pair, such as is returned by tdm_tidiers, containing columns document and term. It returns that same data with an additional column `.topic` with the topic assignment for that document-term combination.

See Also

`lda_tidiers`, `mallet.doc.topics`, `mallet.topic.words`

Examples

```
## Not run:
library(mallet)
library(dplyr)

data("AssociatedPress", package = "topicmodels")
td <- tidy(AssociatedPress)

# mallet needs a file with stop words
tmp <- tempfile()
writeLines(stop_words$word, tmp)

# two vectors: one with document IDs, one with text
docs <- td %>%
  group_by(document = as.character(document)) %>%
  summarize(text = paste(rep(term, count), collapse = " "))

docs <- mallet.import(docs$document, docs$text, tmp)

# create and run a topic model
topic_model <- MalletLDA(num.topics = 4)
topic_model$loadDocuments(docs)
topic_model$train(20)

# tidy the word-topic combinations
td_beta <- tidy(topic_model)
td_beta

# Examine the four topics
td_beta %>%
  group_by(topic) %>%
  top_n(8, beta) %>%
  ungroup() %>%
  mutate(term = reorder(term, beta)) %>%
```

```
    ggplot(aes(term, beta)) +
    geom_col() +
    facet_wrap(~ topic, scales = "free") +
    coord_flip()

# find the assignments of each word in each document
assignments <- augment(topic_model, td)
assignments

## End(Not run)
```

---

nma_words   *English negators, modals, and adverbs*

Description

English negators, modals, and adverbs, as a data frame. A few of these entries are two-word phrases instead of single words.

Usage

```
nma_words
```

Format

A data frame with 44 rows and 2 variables:

word An English word or bigram

modifier The modifier type for word, either "negator", "modal", or "adverb"

Source

<http://saifmohammad.com/WebPages/SCL.html#NMA>

---

parts_of_speech   *Parts of speech for English words from the Moby Project*

Description

Parts of speech for English words from the Moby Project by Grady Ward. Words with non-ASCII characters and items with a space have been removed.

Usage

```
parts_of_speech
```

Format

A data frame with 205,985 rows and 2 variables:

word An English word

pos The part of speech of the word. One of 13 options, such as "Noun", "Adverb", "Adjective"

Details

Another dataset of English parts of speech, available only for non-commercial use, is available as part of SUBTLEXus at https://www.ugent.be/pp/experimentele-psychologie/en/research/documents/subtlexus/.

Source

https://archive.org/details/mobypartofspeech03203gut

Examples

```
library(dplyr)

parts_of_speech

parts_of_speech %>%
  count(pos, sort = TRUE)
```

---

reorder_within   *Reorder an x or y axis within facets*

Description

Reorder a column before plotting with faceting, such that the values are ordered within each facet. This requires two functions: `reorder_within` applied to the column, then either `scale_x_reordered` or `scale_y_reordered` added to the plot. This is implemented as a bit of a hack: it appends ___ and then the facet at the end of each string.

Usage

```
reorder_within(x, by, within, fun = mean, sep = "___", ...)

scale_x_reordered(..., sep = "___")

scale_y_reordered(..., sep = "___")
```

Arguments

`x`	Vector to reorder.
`by`	Vector of the same length, to use for reordering.
`within`	Vector or list of vectors of the same length that will later be used for faceting. A list of vectors will be used to facet within multiple variables.
`fun`	Function to perform within each subset to determine the resulting ordering. By default, mean.
`sep`	Separator to distinguish the two. You may want to set this manually if ___ can exist within one of your labels.
`...`	In `reorder_within` arguments passed on to `reorder`. In the scale functions, extra arguments passed on to `scale_x_discrete` or `scale_y_discrete`.

Source

"Ordering categories within ggplot2 Facets" by Tyler Rinker: `https://trinkerrstuff.wordpress.com/2016/12/23/ordering-categories-within-ggplot2-facets/`

Examples

```
library(tidyr)
library(ggplot2)

iris_gathered <- gather(iris, metric, value, -Species)

# reordering doesn't work within each facet (see Sepal.Width):
ggplot(iris_gathered, aes(reorder(Species, value), value)) +
  geom_boxplot() +
  facet_wrap(~ metric)

# reorder_within and scale_x_reordered work.
# (Note that you need to set scales = "free_x" in the facet)
ggplot(iris_gathered, aes(reorder_within(Species, value, metric), value)) +
  geom_boxplot() +
  scale_x_reordered() +
  facet_wrap(~ metric, scales = "free_x")

# to reorder within multiple variables, set within to the list of
# facet variables.
ggplot(mtcars, aes(reorder_within(carb, mpg, list(vs, am)), mpg)) +
  geom_boxplot() +
  scale_x_reordered() +
  facet_wrap(vs ~ am, scales = "free_x")
```

---

sentiments   *Sentiment lexicon from Bing Liu and collaborators*

Description

Lexicon for opinion and sentiment analysis in a tidy data frame. This dataset is included in this package with permission of the creators, and may be used in research, commercial, etc. contexts with attribution, using either the paper or URL below.

Usage

    sentiments

Format

A data frame with 6,786 rows and 2 variables:

word An English word

sentiment A sentiment for that word, either positive or negative.

Details

This lexicon was first published in:

Minqing Hu and Bing Liu, "Mining and summarizing customer reviews.", Proceedings of the ACM SIGKDD International Conference on Knowledge Discovery & Data Mining (KDD-2004), Seattle, Washington, USA, Aug 22-25, 2004.

Words with non-ASCII characters were removed.

Source

<https://www.cs.uic.edu/~liub/FBS/sentiment-analysis.html>

---

stm_tidiers   *Tidiers for Structural Topic Models from the stm package*

Description

Tidy topic models fit by the stm package. The arguments and return values are similar to `lda_tidiers`.

Usage

```
## S3 method for class 'STM'
tidy(
  x,
  matrix = c("beta", "gamma", "theta"),
  log = FALSE,
  document_names = NULL,
  ...
)

## S3 method for class 'estimateEffect'
tidy(x, ...)

## S3 method for class 'estimateEffect'
glance(x, ...)

## S3 method for class 'STM'
augment(x, data, ...)

## S3 method for class 'STM'
glance(x, ...)
```

Arguments

x	An STM fitted model object from either `stm` or `estimateEffect` from the stm package.
matrix	Whether to tidy the beta (per-term-per-topic, default) or gamma/theta (per-document-per-topic) matrix. The stm package calls this the theta matrix, but other topic modeling packages call this gamma.
log	Whether beta/gamma/theta should be on a log scale, default FALSE
document_names	Optional vector of document names for use with per-document-per-topic tidying
...	Extra arguments, not used
data	For augment, the data given to the stm function, either as a `dfm` from quanteda or as a tidied table with "document" and "term" columns

Value

`tidy` returns a tidied version of either the beta or gamma matrix if called on an object from `stm` or a tidied version of the estimated regressions if called on an object from `estimateEffect`.

`glance` always returns a one-row table, with columns

k Number of topics in the model

docs Number of documents in the model

uncertainty Uncertainty measure

augment must be provided a data argument, either a `dfm` from quanteda or a table containing one row per original document-term pair, such as is returned by [tdm_tidiers](#), containing columns `document` and `term`. It returns that same data as a table with an additional column `.topic` with the topic assignment for that document-term combination.

glance always returns a one-row table, with columns

k Number of topics in the model

docs Number of documents in the model

terms Number of terms in the model

iter Number of iterations used

alpha If an LDA model, the parameter of the Dirichlet distribution for topics over documents

See Also

[lda_tidiers](#)

If `matrix == "beta"` (default), returns a table with one row per topic and term, with columns

topic Topic, as an integer

term Term

beta Probability of a term generated from a topic according to the structural topic model

If `matrix == "gamma"`, returns a table with one row per topic and document, with columns

topic Topic, as an integer

document Document name (if given in vector of `document_names`) or ID as an integer

gamma Probability of topic given document

If called on an object from `estimateEffect`, returns a table with columns

topic Topic, as an integer

term The term in the model being estimated and tested

estimate The estimated coefficient

std.error The standard error from the linear model

statistic t-statistic

p.value two-sided p-value

Examples

```
## Not run:
if (requireNamespace("stm", quietly = TRUE)) {
  library(dplyr)
  library(ggplot2)
  library(stm)
  library(janeaustenr)

  austen_sparse <- austen_books() %>%
```

```
      unnest_tokens(word, text) %>%
      anti_join(stop_words) %>%
      count(book, word) %>%
      cast_sparse(book, word, n)
  topic_model <- stm(austen_sparse, K = 12, verbose = FALSE, init.type = "Spectral")

  # tidy the word-topic combinations
  td_beta <- tidy(topic_model)
  td_beta

  # Examine the topics
  td_beta %>%
    group_by(topic) %>%
    top_n(10, beta) %>%
    ungroup() %>%
    ggplot(aes(term, beta)) +
    geom_col() +
    facet_wrap(~ topic, scales = "free") +
    coord_flip()

  # tidy the document-topic combinations, with optional document names
  td_gamma <- tidy(topic_model, matrix = "gamma",
                   document_names = rownames(austen_sparse))
  td_gamma

  # using stm's gardarianFit, we can tidy the result of a model
  # estimated with covariates
  effects <- estimateEffect(1:3 ~ treatment, gadarianFit, gadarian)
  glance(effects)
  td_estimate <- tidy(effects)
  td_estimate

}

## End(Not run)
```

---

stop_words   *Various lexicons for English stop words*

Description

English stop words from three lexicons, as a data frame. The snowball and SMART sets are pulled
from the tm package. Note that words with non-ASCII characters have been removed.

Usage

```
stop_words
```

Format

A data frame with 1149 rows and 2 variables:

word An English word

lexicon The source of the stop word. Either "onix", "SMART", or "snowball"

Source

- <http://www.lextek.com/manuals/onix/stopwords1.html>
- <https://www.jmlr.org/papers/volume5/lewis04a/lewis04a.pdf>
- <http://snowball.tartarus.org/algorithms/english/stop.txt>

---

tdm_tidiers   *Tidy DocumentTermMatrix, TermDocumentMatrix, and related objects from the tm package*

Description

Tidy a DocumentTermMatrix or TermDocumentMatrix into a three-column data frame: term{}, and value (with zeros missing), with one-row-per-term-per-document.

Usage

```
## S3 method for class 'DocumentTermMatrix'
tidy(x, ...)

## S3 method for class 'TermDocumentMatrix'
tidy(x, ...)

## S3 method for class 'dfm'
tidy(x, ...)

## S3 method for class 'dfmSparse'
tidy(x, ...)

## S3 method for class 'simple_triplet_matrix'
tidy(x, row_names = NULL, col_names = NULL, ...)
```

Arguments

x	A DocumentTermMatrix or TermDocumentMatrix object
...	Extra arguments, not used
row_names	Specify row names
col_names	Specify column names

Examples

```
if (requireNamespace("topicmodels", quietly = TRUE)) {
  data("AssociatedPress", package = "topicmodels")
  AssociatedPress

  tidy(AssociatedPress)
}
```

---

tidy.Corpus	*Tidy a Corpus object from the tm package*

Description

Tidy a Corpus object from the tm package. Returns a data frame with one-row-per-document, with a text column containing the document's text, and one column for each local (per-document) metadata tag. For corpus objects from the quanteda package, see `tidy.corpus`.

Usage

```
## S3 method for class 'Corpus'
tidy(x, collapse = "\n", ...)
```

Arguments

x	A Corpus object, such as a VCorpus or PCorpus
collapse	A string that should be used to collapse text within each corpus (if a document has multiple lines). Give NULL to not collapse strings, in which case a corpus will end up as a list column if there are multi-line documents.
...	Extra arguments, not used

Examples

```
library(dplyr)    # displaying tbl_dfs

if (requireNamespace("tm", quietly = TRUE)) {
  library(tm)
  #' # tm package examples
  txt <- system.file("texts", "txt", package = "tm")
  ovid <- VCorpus(DirSource(txt, encoding = "UTF-8"),
                  readerControl = list(language = "lat"))

  ovid
  tidy(ovid)

  # choose different options for collapsing text within each
```

```
  # document
  tidy(ovid, collapse = "")$text
  tidy(ovid, collapse = NULL)$text

  # another example from Reuters articles
  reut21578 <- system.file("texts", "crude", package = "tm")
  reuters <- VCorpus(DirSource(reut21578),
                     readerControl = list(reader = readReut21578XMLasPlain))
  reuters

  tidy(reuters)
}
```

---

tidytext   *tidytext: Text Mining using 'dplyr', 'ggplot2', and Other Tidy Tools*

Description

This package implements tidy data principles to make many text mining tasks easier, more effective, and consistent with tools already in wide use.

Details

Much of the infrastructure needed for text mining with tidy data frames already exists in packages like dplyr, broom, tidyr and ggplot2.

In this package, we provide functions and supporting data sets to allow conversion of text to and from tidy formats, and to switch seamlessly between tidy tools and existing text mining packages.

To learn more about tidytext, start with the vignettes: `browseVignettes(package = "tidytext")`

---

tidy_triplet   *Utility function to tidy a simple triplet matrix*

Description

Utility function to tidy a simple triplet matrix

Usage

```
tidy_triplet(x, triplets, row_names = NULL, col_names = NULL)
```

Arguments

x	Object with rownames and colnames
triplets	A data frame or list of i, j, x
row_names	rownames, if not gotten from rownames(x)
col_names	colnames, if not gotten from colnames(x)

---

unnest_characters   *Wrapper around unnest_tokens for characters and character shingles*

Description

These functions are a wrapper around unnest_tokens(token = "characters") and unnest_tokens(
token = "character_shingles").

Usage

```
unnest_characters(
  tbl,
  output,
  input,
  strip_non_alphanum = TRUE,
  format = c("text", "man", "latex", "html", "xml"),
  to_lower = TRUE,
  drop = TRUE,
  collapse = NULL,
  ...
)

unnest_character_shingles(
  tbl,
  output,
  input,
  n = 3L,
  n_min = n,
  strip_non_alphanum = TRUE,
  format = c("text", "man", "latex", "html", "xml"),
  to_lower = TRUE,
  drop = TRUE,
  collapse = NULL,
  ...
)
```

Arguments

tbl	A data frame
output	Output column to be created as string or symbol.
input	Input column that gets split as string or symbol.
	The output/input arguments are passed by expression and support quasiquotation; you can unquote strings and symbols.
strip_non_alphanum	
	Should punctuation and white space be stripped?

format	Either "text", "man", "latex", "html", or "xml". When the format is "text", this function uses the tokenizers package. If not "text", this uses the hunspell tokenizer, and can tokenize only by "word"
to_lower	Whether to convert tokens to lowercase. If tokens include URLS (such as with token = "tweets"), such converted URLs may no longer be correct.
drop	Whether original input column should get dropped. Ignored if the original input and new output column have the same name.
collapse	A character vector of variables to collapse text across, or NULL.
	For tokens like n-grams or sentences, text can be collapsed across rows within variables specified by collapse before tokenization. At tidytext 0.2.7, the default behavior for collapse = NULL changed to be more consistent. The new behavior is that text is *not* collapsed for NULL.
	Grouping data specifies variables to collapse across in the same way as collapse but you **cannot** use both the collapse argument and grouped data. Collapsing applies mostly to token options of "ngrams", "skip_ngrams", "sentences", "lines", "paragraphs", or "regex".
...	Extra arguments passed on to [tokenizers](#)
n	The number of characters in each shingle. This must be an integer greater than or equal to 1.
n_min	This must be an integer greater than or equal to 1, and less than or equal to n.

See Also

- [unnest_tokens()](#)

Examples

```
library(dplyr)
library(janeaustenr)

d <- tibble(txt = prideprejudice)

d %>%
  unnest_characters(word, txt)

d %>%
  unnest_character_shingles(word, txt, n = 3)
```

---

unnest_ngrams	*Wrapper around unnest_tokens for n-grams*

Description

These functions are wrappers around unnest_tokens(token = "ngrams") and unnest_tokens(token = "skip_ngrams") .

Usage

```
unnest_ngrams(
  tbl,
  output,
  input,
  n = 3L,
  n_min = n,
  ngram_delim = " ",
  format = c("text", "man", "latex", "html", "xml"),
  to_lower = TRUE,
  drop = TRUE,
  collapse = NULL,
  ...
)

unnest_skip_ngrams(
  tbl,
  output,
  input,
  n_min = 1,
  n = 3,
  k = 1,
  format = c("text", "man", "latex", "html", "xml"),
  to_lower = TRUE,
  drop = TRUE,
  collapse = NULL,
  ...
)
```

Arguments

`tbl`	A data frame
`output`	Output column to be created as string or symbol.
`input`	Input column that gets split as string or symbol.
	The output/input arguments are passed by expression and support [quasiquotation](#); you can unquote strings and symbols.
`n`	The number of words in the n-gram. This must be an integer greater than or equal to 1.
`n_min`	This must be an integer greater than or equal to 1, and less than or equal to n.
`ngram_delim`	The separator between words in an n-gram.
`format`	Either "text", "man", "latex", "html", or "xml". When the format is "text", this function uses the tokenizers package. If not "text", this uses the hunspell tokenizer, and can tokenize only by "word"
`to_lower`	Whether to convert tokens to lowercase. If tokens include URLS (such as with `token = "tweets"`), such converted URLs may no longer be correct.

drop	Whether original input column should get dropped. Ignored if the original input and new output column have the same name.
collapse	A character vector of variables to collapse text across, or NULL.
	For tokens like n-grams or sentences, text can be collapsed across rows within variables specified by collapse before tokenization. At tidytext 0.2.7, the default behavior for collapse = NULL changed to be more consistent. The new behavior is that text is *not* collapsed for NULL.
	Grouping data specifies variables to collapse across in the same way as collapse but you **cannot** use both the collapse argument and grouped data. Collapsing applies mostly to token options of "ngrams", "skip_ngrams", "sentences", "lines", "paragraphs", or "regex".
...	Extra arguments passed on to [tokenizers](#)
k	For the skip n-gram tokenizer, the maximum skip distance between words. The function will compute all skip n-grams between 0 and k.

See Also

- [unnest_tokens()](#)

Examples

```
library(dplyr)
library(janeaustenr)

d <- tibble(txt = prideprejudice)

d %>%
  unnest_ngrams(word, txt, n = 2)

d %>%
  unnest_skip_ngrams(word, txt, n = 3, k = 1)
```

---

unnest_ptb   *Wrapper around unnest_tokens for Penn Treebank Tokenizer*

Description

This function is a wrapper around unnest_tokens(token = "ptb").

Usage

```
unnest_ptb(
  tbl,
  output,
  input,
  format = c("text", "man", "latex", "html", "xml"),
```

```
    to_lower = TRUE,
    drop = TRUE,
    collapse = NULL,
    ...
)
```

Arguments

`tbl`	A data frame
`output`	Output column to be created as string or symbol.
`input`	Input column that gets split as string or symbol.
	The output/input arguments are passed by expression and support [quasiquotation](); you can unquote strings and symbols.
`format`	Either "text", "man", "latex", "html", or "xml". When the format is "text", this function uses the tokenizers package. If not "text", this uses the hunspell tokenizer, and can tokenize only by "word"
`to_lower`	Whether to convert tokens to lowercase. If tokens include URLS (such as with `token = "tweets"`), such converted URLs may no longer be correct.
`drop`	Whether original input column should get dropped. Ignored if the original input and new output column have the same name.
`collapse`	A character vector of variables to collapse text across, or `NULL`.
	For tokens like n-grams or sentences, text can be collapsed across rows within variables specified by `collapse` before tokenization. At tidytext 0.2.7, the default behavior for `collapse = NULL` changed to be more consistent. The new behavior is that text is *not* collapsed for `NULL`.
	Grouping data specifies variables to collapse across in the same way as `collapse` but you **cannot** use both the `collapse` argument and grouped data. Collapsing applies mostly to token options of "ngrams", "skip_ngrams", "sentences", "lines", "paragraphs", or "regex".
`...`	Extra arguments passed on to [tokenizers]()

See Also

- [unnest_tokens()]()

Examples

```
library(dplyr)
library(janeaustenr)

d <- tibble(txt = prideprejudice)

d %>%
  unnest_ptb(word, txt)
```

unnest_regex	*Wrapper around unnest_tokens for regular expressions*

Description

This function is a wrapper around unnest_tokens(token = ″regex″).

Usage

```
unnest_regex(
  tbl,
  output,
  input,
  pattern = ″\\s+″,
  format = c(″text″, ″man″, ″latex″, ″html″, ″xml″),
  to_lower = TRUE,
  drop = TRUE,
  collapse = NULL,
  ...
)
```

Arguments

tbl	A data frame
output	Output column to be created as string or symbol.
input	Input column that gets split as string or symbol.
	The output/input arguments are passed by expression and support [quasiquotation](#); you can unquote strings and symbols.
pattern	A regular expression that defines the split.
format	Either "text", "man", "latex", "html", or "xml". When the format is "text", this function uses the tokenizers package. If not "text", this uses the hunspell tokenizer, and can tokenize only by "word"
to_lower	Whether to convert tokens to lowercase. If tokens include URLS (such as with token = ″tweets″), such converted URLs may no longer be correct.
drop	Whether original input column should get dropped. Ignored if the original input and new output column have the same name.
collapse	A character vector of variables to collapse text across, or NULL.
	For tokens like n-grams or sentences, text can be collapsed across rows within variables specified by collapse before tokenization. At tidytext 0.2.7, the default behavior for collapse = NULL changed to be more consistent. The new behavior is that text is *not* collapsed for NULL.
	Grouping data specifies variables to collapse across in the same way as collapse but you **cannot** use both the collapse argument and grouped data. Collapsing applies mostly to token options of "ngrams", "skip_ngrams", "sentences", "lines", "paragraphs", or "regex".
...	Extra arguments passed on to [tokenizers](#)

See Also

- unnest_tokens()

Examples

```
library(dplyr)
library(janeaustenr)

d <- tibble(txt = prideprejudice)

d %>%
  unnest_regex(word, txt, pattern = "Chapter [\\\\d]")
```

---

unnest_sentences   *Wrapper around unnest_tokens for sentences, lines, and paragraphs*

Description

These functions are wrappers around unnest_tokens(token = "sentences") unnest_tokens(token = "lines") and unnest_tokens(token = "paragraphs").

Usage

```
unnest_sentences(
  tbl,
  output,
  input,
  strip_punct = FALSE,
  format = c("text", "man", "latex", "html", "xml"),
  to_lower = TRUE,
  drop = TRUE,
  collapse = NULL,
  ...
)

unnest_lines(
  tbl,
  output,
  input,
  format = c("text", "man", "latex", "html", "xml"),
  to_lower = TRUE,
  drop = TRUE,
  collapse = NULL,
  ...
)
```

```
unnest_paragraphs(
  tbl,
  output,
  input,
  paragraph_break = "\n\n",
  format = c("text", "man", "latex", "html", "xml"),
  to_lower = TRUE,
  drop = TRUE,
  collapse = NULL,
  ...
)
```

Arguments

`tbl`	A data frame
`output`	Output column to be created as string or symbol.
`input`	Input column that gets split as string or symbol.
	The output/input arguments are passed by expression and support [quasiquotation](); you can unquote strings and symbols.
`strip_punct`	Should punctuation be stripped?
`format`	Either "text", "man", "latex", "html", or "xml". When the format is "text", this function uses the tokenizers package. If not "text", this uses the hunspell tokenizer, and can tokenize only by "word"
`to_lower`	Whether to convert tokens to lowercase. If tokens include URLS (such as with `token = "tweets"`), such converted URLs may no longer be correct.
`drop`	Whether original input column should get dropped. Ignored if the original input and new output column have the same name.
`collapse`	A character vector of variables to collapse text across, or `NULL`.
	For tokens like n-grams or sentences, text can be collapsed across rows within variables specified by `collapse` before tokenization. At tidytext 0.2.7, the default behavior for `collapse = NULL` changed to be more consistent. The new behavior is that text is *not* collapsed for `NULL`.
	Grouping data specifies variables to collapse across in the same way as `collapse` but you **cannot** use both the `collapse` argument and grouped data. Collapsing applies mostly to token options of "ngrams", "skip_ngrams", "sentences", "lines", "paragraphs", or "regex".
`...`	Extra arguments passed on to [tokenizers]()
`paragraph_break`	
	A string identifying the boundary between two paragraphs.

See Also

- [unnest_tokens()]()

Examples

```
library(dplyr)
library(janeaustenr)

d <- tibble(txt = prideprejudice)

d %>%
  unnest_sentences(word, txt)
```

---

unnest_tokens *Split a column into tokens*

Description

Split a column into tokens, flattening the table into one-token-per-row. This function supports non-standard evaluation through the tidyeval framework.

Usage

```
unnest_tokens(
  tbl,
  output,
  input,
  token = "words",
  format = c("text", "man", "latex", "html", "xml"),
  to_lower = TRUE,
  drop = TRUE,
  collapse = NULL,
  ...
)
```

Arguments

tbl	A data frame
output	Output column to be created as string or symbol.
input	Input column that gets split as string or symbol.
	The output/input arguments are passed by expression and support quasiquotation; you can unquote strings and symbols.
token	Unit for tokenizing, or a custom tokenizing function. Built-in options are "words" (default), "characters", "character_shingles", "ngrams", "skip_ngrams", "sentences", "lines", "paragraphs", "regex", "tweets" (tokenization by word that preserves usernames, hashtags, and URLS), and "ptb" (Penn Treebank). If a function, should take a character vector and return a list of character vectors of the same length.

format	Either "text", "man", "latex", "html", or "xml". When the format is "text", this function uses the tokenizers package. If not "text", this uses the hunspell tokenizer, and can tokenize only by "word"
to_lower	Whether to convert tokens to lowercase. If tokens include URLS (such as with token = "tweets"), such converted URLs may no longer be correct.
drop	Whether original input column should get dropped. Ignored if the original input and new output column have the same name.
collapse	A character vector of variables to collapse text across, or NULL.
	For tokens like n-grams or sentences, text can be collapsed across rows within variables specified by collapse before tokenization. At tidytext 0.2.7, the default behavior for collapse = NULL changed to be more consistent. The new behavior is that text is *not* collapsed for NULL.
	Grouping data specifies variables to collapse across in the same way as collapse but you **cannot** use both the collapse argument and grouped data. Collapsing applies mostly to token options of "ngrams", "skip_ngrams", "sentences", "lines", "paragraphs", or "regex".
...	Extra arguments passed on to [tokenizers](), such as strip_punct for "words" and "tweets", n and k for "ngrams" and "skip_ngrams", strip_url for "tweets", and pattern for "regex".

Details

If format is anything other than "text", this uses the [hunspell_parse]() tokenizer instead of the tokenizers package. This does not yet have support for tokenizing by any unit other than words.

Examples

```
library(dplyr)
library(janeaustenr)

d <- tibble(txt = prideprejudice)
d

d %>%
  unnest_tokens(word, txt)

d %>%
  unnest_tokens(sentence, txt, token = "sentences")

d %>%
  unnest_tokens(ngram, txt, token = "ngrams", n = 2)

d %>%
  unnest_tokens(chapter, txt, token = "regex", pattern = "Chapter [\\\\d]")

d %>%
  unnest_tokens(shingle, txt, token = "character_shingles", n = 4)
```

```
# custom function
d %>%
  unnest_tokens(word, txt, token = stringr::str_split, pattern = " ")

# tokenize HTML
h <- tibble(row = 1:2,
              text = c("<h1>Text <b>is</b>", "<a href='example.com'>here</a>"))

h %>%
  unnest_tokens(word, text, format = "html")
```

---

unnest_tweets   *Wrapper around unnest_tokens for tweets*

Description

This function is a wrapper around unnest_tokens(token = "tweets").

Usage

```
unnest_tweets(
  tbl,
  output,
  input,
  strip_punct = TRUE,
  strip_url = FALSE,
  format = c("text", "man", "latex", "html", "xml"),
  to_lower = TRUE,
  drop = TRUE,
  collapse = NULL,
  ...
)
```

Arguments

tbl	A data frame
output	Output column to be created as string or symbol.
input	Input column that gets split as string or symbol.
	The output/input arguments are passed by expression and support [quasiquotation](); you can unquote strings and symbols.
strip_punct	Should punctuation be stripped?
strip_url	Should URLs (starting with http(s)) be preserved intact, or removed entirely?
format	Either "text", "man", "latex", "html", or "xml". When the format is "text", this function uses the tokenizers package. If not "text", this uses the hunspell tokenizer, and can tokenize only by "word"

to_lower	Whether to convert tokens to lowercase. If tokens include URLS (such as with token = "tweets"), such converted URLs may no longer be correct.
drop	Whether original input column should get dropped. Ignored if the original input and new output column have the same name.
collapse	A character vector of variables to collapse text across, or NULL.
	For tokens like n-grams or sentences, text can be collapsed across rows within variables specified by collapse before tokenization. At tidytext 0.2.7, the default behavior for collapse = NULL changed to be more consistent. The new behavior is that text is *not* collapsed for NULL.
	Grouping data specifies variables to collapse across in the same way as collapse but you **cannot** use both the collapse argument and grouped data. Collapsing applies mostly to token options of "ngrams", "skip_ngrams", "sentences", "lines", "paragraphs", or "regex".
...	Extra arguments passed on to tokenizers

See Also

- unnest_tokens()

Examples

```
library(dplyr)
tweets <- tibble(
   id = 1,
   txt = "@rOpenSci and #rstats see: https://cran.r-project.org"
)

tweets %>%
   unnest_tweets(out, txt)
```

Index